

# Capstone\_Project(2)

September 7, 2020

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (you could download the notebook with File -> Download .ipynb, open the notebook locally, and then File -> Download as -> PDF via LaTeX), and then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
[82]: import tensorflow as tf
      from scipy.io import loadmat
      import numpy as np
      import matplotlib.pyplot as plt
      from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
      from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,   
↳Dropout, BatchNormalization  
import random
```

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

The train and test datasets required for this project can be downloaded from [here](#) and [here](#). Once unzipped, you will have two files: `train_32x32.mat` and `test_32x32.mat`. You should store these files in Drive for use in this Colab notebook.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

[5]: *# Run this cell to connect to your Drive folder*

```
from google.colab import drive  
drive.mount('/content/gdrive')
```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly&response\\_type=code](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly&response_type=code)

Enter your authorization code:

.....

Mounted at /content/gdrive

[6]: *# Load the dataset from your Drive folder*

```
train = loadmat('gdrive/My Drive/train_32x32.mat')  
test = loadmat('gdrive/My Drive/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.

- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
[7]: # Loading the dataset
```

```
x_train = train['X']
x_test = test['X']
y_train = train['y']
y_test = test['y']
```

```
[8]: # analysing the dimensions of the input
x_train.shape, x_test.shape
```

```
[8]: ((32, 32, 3, 73257), (32, 32, 3, 26032))
```

```
[9]: # changing dimensions from (a, b, c, num_examples) to (num_examples, a, b, c)
```

```
x_train = np.moveaxis(x_train, -1, 0)
x_test = np.moveaxis(x_test, -1, 0)
```


```
[10]: # checking if dimensions have actually changed
```

```
x_train.shape, x_test.shape
```


```
[10]: ((73257, 32, 32, 3), (26032, 32, 32, 3))
```

```
[11]: # plotting some sample images
```


```
for i in range(10):
    plt.imshow(x_train[i, :, :, :])
    plt.show()
    print(y_train[i])
```




[1]




[9]




[2]




[3]




[2]




[5]




[9]



[3]



[3]



[1]




```
[12]: # making changes to the images
```


```
x_train_gray = np.mean(x_train, 3).reshape(73257, 32, 32, 1) / 255.  
x_test_gray = np.mean(x_test, 3).reshape(26032, 32, 32, 1) / 255.  
x_train_plot = np.mean(x_train, 3)
```

```
[13]: # plotting the training images
```


```
for i in range(10):  
    plt.imshow(x_train_plot[i, :, :,])  
    plt.show()  
    print(y_train[i])
```




```
[1]
```




[9]




[2]




[3]




[2]




[5]




[9]



[3]



[3]



[1]

```
[62]: # was having trouble in MLP NN classifier, so converting to one-hot labels
```

```
x_train[0].shape
```

```
[62]: (32, 32, 3)
```

```
[63]: from sklearn.preprocessing import OneHotEncoder
```


```
enc = OneHotEncoder().fit(y_train)
y_train_oh = enc.transform(y_train).toarray()
y_test_oh = enc.transform(y_test).toarray()
```

```
[64]: y_test_oh[0]
```

```
[64]: array([0., 0., 0., 0., 1., 0., 0., 0., 0., 0.])
```

```
[65]: plt.imshow(x_test[0])
```

```
[65]: <matplotlib.image.AxesImage at 0x7fd52a5facf8>
```



### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[18]: model_seq = Sequential([
        Flatten(input_shape=x_train[0].shape),
        Dense(128, activation='relu'),
        Dense(256, activation='relu'),
        BatchNormalization(),
        Dense(256, activation='relu'),
        Dropout(0.5),
```

```

        Dense(512, activation='relu'),
        Dense(10, activation='softmax')

])

model_seq.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 128)	393344
dense_1 (Dense)	(None, 256)	33024
batch_normalization (Batch Normalization)	(None, 256)	1024
dense_2 (Dense)	(None, 256)	65792
dropout (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 512)	131584
dense_4 (Dense)	(None, 10)	5130

Total params: 629,898  
 Trainable params: 629,386  
 Non-trainable params: 512

```

[19]: model_seq.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

```

```

[20]: checkpoint = ModelCheckpoint(filepath='sequential',
                                   save_best_only=True,
                                   save_weights_only=True,
                                   monitor='val_loss',
                                   mode='min',
                                   verbose=1)

early_stop = EarlyStopping(patience=5, monitor='loss')

```

```

[21]: history = model_seq.fit(x_train, y_train_oh, epochs=30,
                              validation_data=(x_test, y_test_oh),

```



```
batch_size=128,  
callbacks=[checkpoint, early_stop])
```

```
Epoch 1/30  
564/573 [=====>.] - ETA: 0s - loss: 1.7239 - accuracy:  
0.3994  
Epoch 00001: val_loss improved from inf to 2.01975, saving model to sequential  
573/573 [=====] - 3s 4ms/step - loss: 1.7188 -  
accuracy: 0.4018 - val_loss: 2.0197 - val_accuracy: 0.3412  
Epoch 2/30  
559/573 [=====>.] - ETA: 0s - loss: 1.2283 - accuracy:  
0.6102  
Epoch 00002: val_loss did not improve from 2.01975  
573/573 [=====] - 2s 4ms/step - loss: 1.2277 -  
accuracy: 0.6104 - val_loss: 2.3677 - val_accuracy: 0.3506  
Epoch 3/30  
563/573 [=====>.] - ETA: 0s - loss: 1.1261 - accuracy:  
0.6475  
Epoch 00003: val_loss improved from 2.01975 to 1.91625, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 1.1250 -  
accuracy: 0.6476 - val_loss: 1.9162 - val_accuracy: 0.4145  
Epoch 4/30  
566/573 [=====>.] - ETA: 0s - loss: 1.0364 - accuracy:  
0.6775  
Epoch 00004: val_loss improved from 1.91625 to 1.54374, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 1.0353 -  
accuracy: 0.6778 - val_loss: 1.5437 - val_accuracy: 0.5225  
Epoch 5/30  
571/573 [=====>.] - ETA: 0s - loss: 0.9973 - accuracy:  
0.6885  
Epoch 00005: val_loss did not improve from 1.54374  
573/573 [=====] - 2s 4ms/step - loss: 0.9970 -  
accuracy: 0.6886 - val_loss: 1.8820 - val_accuracy: 0.4585  
Epoch 6/30  
572/573 [=====>.] - ETA: 0s - loss: 0.9263 - accuracy:  
0.7130  
Epoch 00006: val_loss improved from 1.54374 to 1.14499, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.9263 -  
accuracy: 0.7130 - val_loss: 1.1450 - val_accuracy: 0.6452  
Epoch 7/30  
566/573 [=====>.] - ETA: 0s - loss: 0.8867 - accuracy:  
0.7239  
Epoch 00007: val_loss improved from 1.14499 to 1.03599, saving model to  
sequential
```

573/573 [=====] - 2s 4ms/step - loss: 0.8859 - accuracy: 0.7240 - val\_loss: 1.0360 - val\_accuracy: 0.6787  
Epoch 8/30  
556/573 [=====>.] - ETA: 0s - loss: 0.8566 - accuracy: 0.7325  
Epoch 00008: val\_loss improved from 1.03599 to 0.93794, saving model to sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.8560 - accuracy: 0.7329 - val\_loss: 0.9379 - val\_accuracy: 0.7068  
Epoch 9/30  
567/573 [=====>.] - ETA: 0s - loss: 0.8503 - accuracy: 0.7349  
Epoch 00009: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.8508 - accuracy: 0.7350 - val\_loss: 1.1038 - val\_accuracy: 0.6475  
Epoch 10/30  
573/573 [=====] - ETA: 0s - loss: 0.8215 - accuracy: 0.7462  
Epoch 00010: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.8215 - accuracy: 0.7462 - val\_loss: 1.0235 - val\_accuracy: 0.6829  
Epoch 11/30  
573/573 [=====] - ETA: 0s - loss: 0.7973 - accuracy: 0.7530  
Epoch 00011: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7973 - accuracy: 0.7530 - val\_loss: 1.1812 - val\_accuracy: 0.6357  
Epoch 12/30  
568/573 [=====>.] - ETA: 0s - loss: 0.7838 - accuracy: 0.7578  
Epoch 00012: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7839 - accuracy: 0.7578 - val\_loss: 1.0245 - val\_accuracy: 0.6788  
Epoch 13/30  
568/573 [=====>.] - ETA: 0s - loss: 0.7779 - accuracy: 0.7590  
Epoch 00013: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7782 - accuracy: 0.7588 - val\_loss: 1.1422 - val\_accuracy: 0.6450  
Epoch 14/30  
563/573 [=====>.] - ETA: 0s - loss: 0.7591 - accuracy: 0.7627  
Epoch 00014: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7595 - accuracy: 0.7625 - val\_loss: 0.9860 - val\_accuracy: 0.6989  
Epoch 15/30  
557/573 [=====>.] - ETA: 0s - loss: 0.7573 - accuracy: 0.7651


Epoch 00015: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7574 -  
accuracy: 0.7652 - val\_loss: 1.1555 - val\_accuracy: 0.6411  
Epoch 16/30  
564/573 [=====>.] - ETA: 0s - loss: 0.7371 - accuracy:  
0.7715  
Epoch 00016: val\_loss did not improve from 0.93794  
573/573 [=====] - 2s 4ms/step - loss: 0.7361 -  
accuracy: 0.7719 - val\_loss: 1.0144 - val\_accuracy: 0.6901  
Epoch 17/30  
572/573 [=====>.] - ETA: 0s - loss: 0.7236 - accuracy:  
0.7757  
Epoch 00017: val\_loss improved from 0.93794 to 0.92079, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.7236 -  
accuracy: 0.7757 - val\_loss: 0.9208 - val\_accuracy: 0.7214  
Epoch 18/30  
571/573 [=====>.] - ETA: 0s - loss: 0.7132 - accuracy:  
0.7778  
Epoch 00018: val\_loss did not improve from 0.92079  
573/573 [=====] - 2s 4ms/step - loss: 0.7131 -  
accuracy: 0.7779 - val\_loss: 0.9502 - val\_accuracy: 0.7121  
Epoch 19/30  
564/573 [=====>.] - ETA: 0s - loss: 0.7043 - accuracy:  
0.7798  
Epoch 00019: val\_loss improved from 0.92079 to 0.88748, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.7043 -  
accuracy: 0.7798 - val\_loss: 0.8875 - val\_accuracy: 0.7436  
Epoch 20/30  
561/573 [=====>.] - ETA: 0s - loss: 0.6890 - accuracy:  
0.7860  
Epoch 00020: val\_loss did not improve from 0.88748  
573/573 [=====] - 2s 4ms/step - loss: 0.6892 -  
accuracy: 0.7859 - val\_loss: 0.9715 - val\_accuracy: 0.7039  
Epoch 21/30  
564/573 [=====>.] - ETA: 0s - loss: 0.6830 - accuracy:  
0.7866  
Epoch 00021: val\_loss did not improve from 0.88748  
573/573 [=====] - 2s 4ms/step - loss: 0.6820 -  
accuracy: 0.7868 - val\_loss: 0.9741 - val\_accuracy: 0.6929  
Epoch 22/30  
567/573 [=====>.] - ETA: 0s - loss: 0.6806 - accuracy:  
0.7867  
Epoch 00022: val\_loss improved from 0.88748 to 0.84766, saving model to  
sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.6804 -  
accuracy: 0.7868 - val\_loss: 0.8477 - val\_accuracy: 0.7499

Epoch 23/30  
571/573 [=====>.] - ETA: 0s - loss: 0.7068 - accuracy: 0.7807  
Epoch 00023: val\_loss did not improve from 0.84766  
573/573 [=====] - 2s 4ms/step - loss: 0.7064 - accuracy: 0.7808 - val\_loss: 1.0679 - val\_accuracy: 0.6523  
Epoch 24/30  
569/573 [=====>.] - ETA: 0s - loss: 0.6704 - accuracy: 0.7915  
Epoch 00024: val\_loss improved from 0.84766 to 0.83966, saving model to sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.6699 - accuracy: 0.7916 - val\_loss: 0.8397 - val\_accuracy: 0.7399  
Epoch 25/30  
556/573 [=====>.] - ETA: 0s - loss: 0.6583 - accuracy: 0.7944  
Epoch 00025: val\_loss did not improve from 0.83966  
573/573 [=====] - 2s 4ms/step - loss: 0.6580 - accuracy: 0.7946 - val\_loss: 0.8457 - val\_accuracy: 0.7429  
Epoch 26/30  
572/573 [=====>.] - ETA: 0s - loss: 0.6501 - accuracy: 0.7959  
Epoch 00026: val\_loss improved from 0.83966 to 0.81609, saving model to sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.6501 - accuracy: 0.7959 - val\_loss: 0.8161 - val\_accuracy: 0.7528  
Epoch 27/30  
572/573 [=====>.] - ETA: 0s - loss: 0.6621 - accuracy: 0.7928  
Epoch 00027: val\_loss did not improve from 0.81609  
573/573 [=====] - 2s 4ms/step - loss: 0.6620 - accuracy: 0.7928 - val\_loss: 0.9051 - val\_accuracy: 0.7197  
Epoch 28/30  
564/573 [=====>.] - ETA: 0s - loss: 0.6395 - accuracy: 0.7999  
Epoch 00028: val\_loss improved from 0.81609 to 0.76925, saving model to sequential  
573/573 [=====] - 2s 4ms/step - loss: 0.6393 - accuracy: 0.8000 - val\_loss: 0.7692 - val\_accuracy: 0.7642  
Epoch 29/30  
573/573 [=====] - ETA: 0s - loss: 0.6346 - accuracy: 0.8019  
Epoch 00029: val\_loss did not improve from 0.76925  
573/573 [=====] - 2s 4ms/step - loss: 0.6346 - accuracy: 0.8019 - val\_loss: 0.9571 - val\_accuracy: 0.7036  
Epoch 30/30  
565/573 [=====>.] - ETA: 0s - loss: 0.6291 - accuracy: 0.8034

Epoch 00030: val\_loss did not improve from 0.76925  
573/573 [=====] - 2s 4ms/step - loss: 0.6291 -  
accuracy: 0.8033 - val\_loss: 0.9251 - val\_accuracy: 0.7048


```
[24]: plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend(['loss', 'val_loss'], loc='upper right')  
plt.title('Loss')
```

```
[24]: Text(0.5, 1.0, 'Loss')
```



```
[25]: plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.xlabel('Epochs')  
plt.ylabel('Accuracy')  
plt.legend(['accuracy', 'val_accuracy'], loc='upper right')  
plt.title('Accuracy')
```

```
[25]: Text(0.5, 1.0, 'Accuracy')
```



```
[26]: model_seq.evaluate(x_test, y_test_oh, verbose=2)
```

```
814/814 - 1s - loss: 0.9251 - accuracy: 0.7048
```

```
[26]: [0.9251210689544678, 0.7048248052597046]
```

### 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[73]: model_cnn = Sequential([
        Conv2D(16, (3, 3), padding='same', activation='relu',
        ↪input_shape=x_train[0].shape),
        MaxPooling2D((3, 3)),
        Conv2D(32, (3, 3), padding='same', activation='relu'),
        MaxPooling2D((3, 3)),
        BatchNormalization(),
        Conv2D(64, (3, 3), padding='same', activation='relu'),
        MaxPooling2D((3, 3)),
        Dropout(0.5),
        Flatten(),
        Dense(64, activation='relu'),
        Dropout(0.5),
        Dense(10, activation='softmax')
    ])

model_cnn.summary()
```

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
conv2d_33 (Conv2D)	(None, 32, 32, 16)	448
max_pooling2d_33 (MaxPooling)	(None, 10, 10, 16)	0
conv2d_34 (Conv2D)	(None, 10, 10, 32)	4640
max_pooling2d_34 (MaxPooling)	(None, 3, 3, 32)	0
batch_normalization_12 (Batch Normalization)	(None, 3, 3, 32)	128
conv2d_35 (Conv2D)	(None, 3, 3, 64)	18496
max_pooling2d_35 (MaxPooling)	(None, 1, 1, 64)	0
dropout_17 (Dropout)	(None, 1, 1, 64)	0
flatten_12 (Flatten)	(None, 64)	0
dense_30 (Dense)	(None, 64)	4160
dropout_18 (Dropout)	(None, 64)	0
dense_31 (Dense)	(None, 10)	650

=====  
Total params: 28,522

Trainable params: 28,458  
Non-trainable params: 64

-----

```
[74]: model_cnn.compile(optimizer='adam',  
                      loss='categorical_crossentropy',  
                      metrics=['accuracy'])  
  
[75]: checkpoint_cnn = ModelCheckpoint(filepath='CNN', save_best_only=True,  
                                       save_weights_only=True,  
                                       save_freq=5000,  
                                       monitor='val_acc',  
                                       mode='max')  
early_stop_cnn = EarlyStopping(monitor='loss', patience=7, verbose=1)  
  
[76]: history = model_cnn.fit(x_train, y_train_oh,  
                             callbacks=[checkpoint_cnn, early_stop_cnn],  
                             batch_size=128, validation_data=(x_test, y_test_oh),  
                             epochs=30)  
-
```

Epoch 1/30

573/573 [=====] - 3s 5ms/step - loss: 1.9691 -  
accuracy: 0.3054 - val\_loss: 1.2979 - val\_accuracy: 0.5999

Epoch 2/30

573/573 [=====] - 2s 4ms/step - loss: 1.1966 -  
accuracy: 0.6077 - val\_loss: 0.9446 - val\_accuracy: 0.7121

Epoch 3/30

573/573 [=====] - 3s 5ms/step - loss: 0.9402 -  
accuracy: 0.7050 - val\_loss: 0.7039 - val\_accuracy: 0.7875

Epoch 4/30

573/573 [=====] - 3s 5ms/step - loss: 0.8278 -  
accuracy: 0.7458 - val\_loss: 0.6627 - val\_accuracy: 0.8019

Epoch 5/30

573/573 [=====] - 3s 5ms/step - loss: 0.7609 -  
accuracy: 0.7723 - val\_loss: 0.6224 - val\_accuracy: 0.8142

Epoch 6/30

573/573 [=====] - 3s 5ms/step - loss: 0.7127 -  
accuracy: 0.7867 - val\_loss: 0.5515 - val\_accuracy: 0.8375

Epoch 7/30

573/573 [=====] - 2s 4ms/step - loss: 0.6852 -  
accuracy: 0.7956 - val\_loss: 0.5291 - val\_accuracy: 0.8436

Epoch 8/30

573/573 [=====] - 2s 4ms/step - loss: 0.6629 -  
accuracy: 0.8041 - val\_loss: 0.6309 - val\_accuracy: 0.8085

Epoch 9/30

403/573 [=====>...] - ETA: 0s - loss: 0.6416 - accuracy:  
0.8102WARNING:tensorflow:Can save best model only with val\_acc available,



```

skipping.
573/573 [=====] - 2s 4ms/step - loss: 0.6414 -
accuracy: 0.8103 - val_loss: 0.5345 - val_accuracy: 0.8418
Epoch 10/30
573/573 [=====] - 2s 4ms/step - loss: 0.6269 -
accuracy: 0.8154 - val_loss: 0.5266 - val_accuracy: 0.8442
Epoch 11/30
573/573 [=====] - 2s 4ms/step - loss: 0.6155 -
accuracy: 0.8200 - val_loss: 0.6170 - val_accuracy: 0.8127
Epoch 12/30
573/573 [=====] - 2s 4ms/step - loss: 0.6030 -
accuracy: 0.8223 - val_loss: 0.4806 - val_accuracy: 0.8590
Epoch 13/30
573/573 [=====] - 2s 4ms/step - loss: 0.5960 -
accuracy: 0.8255 - val_loss: 0.4718 - val_accuracy: 0.8628
Epoch 14/30
573/573 [=====] - 2s 4ms/step - loss: 0.5871 -
accuracy: 0.8270 - val_loss: 0.4560 - val_accuracy: 0.8651
Epoch 15/30
573/573 [=====] - 2s 4ms/step - loss: 0.5777 -
accuracy: 0.8311 - val_loss: 0.4655 - val_accuracy: 0.8646
Epoch 16/30
573/573 [=====] - 2s 4ms/step - loss: 0.5722 -
accuracy: 0.8321 - val_loss: 0.4537 - val_accuracy: 0.8672
Epoch 17/30
573/573 [=====] - 2s 4ms/step - loss: 0.5654 -
accuracy: 0.8350 - val_loss: 0.4792 - val_accuracy: 0.8594
Epoch 18/30
249/573 [=====>...] - ETA: 1s - loss: 0.5435 - accuracy:
0.8398WARNING:tensorflow:Can save best model only with val_acc available,
skipping.
573/573 [=====] - 2s 4ms/step - loss: 0.5580 -
accuracy: 0.8372 - val_loss: 0.4949 - val_accuracy: 0.8526
Epoch 19/30
573/573 [=====] - 2s 4ms/step - loss: 0.5567 -
accuracy: 0.8385 - val_loss: 0.4576 - val_accuracy: 0.8679
Epoch 20/30
573/573 [=====] - 2s 4ms/step - loss: 0.5483 -
accuracy: 0.8399 - val_loss: 0.4743 - val_accuracy: 0.8601
Epoch 21/30
573/573 [=====] - 2s 4ms/step - loss: 0.5369 -
accuracy: 0.8440 - val_loss: 0.4456 - val_accuracy: 0.8690
Epoch 22/30
573/573 [=====] - 2s 4ms/step - loss: 0.5407 -
accuracy: 0.8429 - val_loss: 0.4500 - val_accuracy: 0.8705
Epoch 23/30
573/573 [=====] - 2s 4ms/step - loss: 0.5390 -
accuracy: 0.8436 - val_loss: 0.4659 - val_accuracy: 0.8636

```

```

Epoch 24/30
573/573 [=====] - 2s 4ms/step - loss: 0.5305 -
accuracy: 0.8463 - val_loss: 0.4548 - val_accuracy: 0.8666
Epoch 25/30
573/573 [=====] - 2s 4ms/step - loss: 0.5304 -
accuracy: 0.8452 - val_loss: 0.4417 - val_accuracy: 0.8705
Epoch 26/30
573/573 [=====] - 2s 4ms/step - loss: 0.5243 -
accuracy: 0.8473 - val_loss: 0.4289 - val_accuracy: 0.8767
Epoch 27/30
 94/573 [==>...] - ETA: 1s - loss: 0.5066 - accuracy:
0.8492WARNING:tensorflow:Can save best model only with val_acc available,
skipping.
573/573 [=====] - 2s 4ms/step - loss: 0.5223 -
accuracy: 0.8465 - val_loss: 0.4881 - val_accuracy: 0.8548
Epoch 28/30
573/573 [=====] - 2s 4ms/step - loss: 0.5253 -
accuracy: 0.8470 - val_loss: 0.4457 - val_accuracy: 0.8685
Epoch 29/30
573/573 [=====] - 2s 4ms/step - loss: 0.5166 -
accuracy: 0.8496 - val_loss: 0.4490 - val_accuracy: 0.8693
Epoch 30/30
573/573 [=====] - 2s 4ms/step - loss: 0.5164 -
accuracy: 0.8488 - val_loss: 0.4219 - val_accuracy: 0.8770

```


[76]: (73257, 10)

```

[78]: plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend(['loss', 'val_loss'], loc='upper right')
      plt.title('Loss')


```

[78]: Text(0.5, 1.0, 'Loss')



```
[80]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['accuracy', 'val_accuracy'], loc='lower right')
plt.title('Accuracy')
```

```
[80]: Text(0.5, 1.0, 'Accuracy')
```



```
[87]: model_cnn.evaluate(x_test, y_test_oh, verbose=2)
```

```
814/814 - 2s - loss: 0.4219 - accuracy: 0.8770
```

```
[87]: [0.4219440221786499, 0.8770359754562378]
```

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
[81]: model_seq.load_weights('sequential')
```

```
[81]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at
0x7fd5dfe8f208>
```

```
[85]: num_test_images = x_test.shape[0]

random_inx = np.random.choice(num_test_images, 5)
random_test_images = x_test[random_inx, ...]
```

```

random_test_labels = y_test[random_inx, ...]

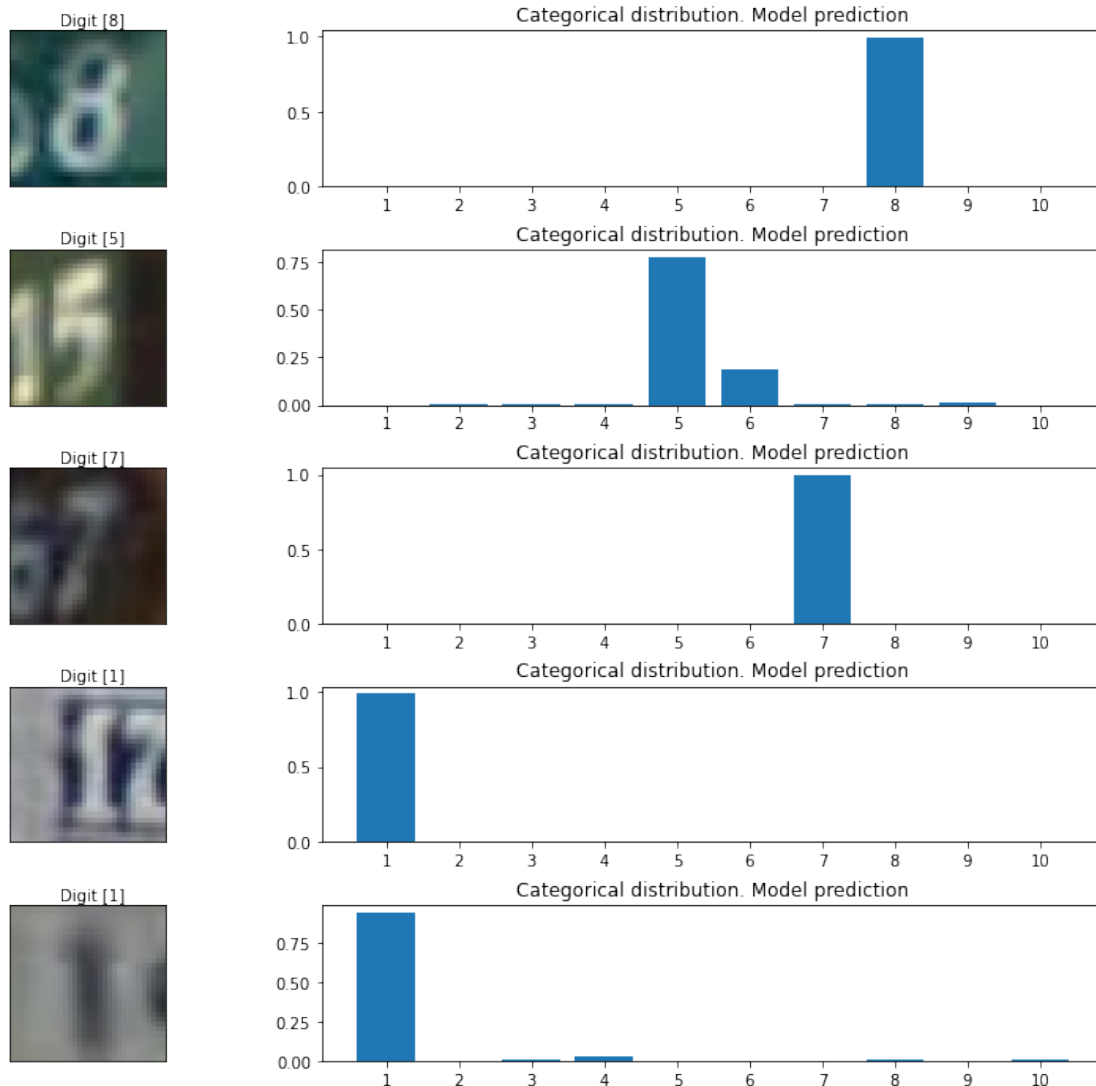
predictions = model_seq.predict(random_test_images)

fig, axes = plt.subplots(5, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)

for i, (prediction, image, label) in enumerate(zip(predictions,
↪random_test_images, random_test_labels)):
    axes[i, 0].imshow(np.squeeze(image))
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Digit {label}')
    axes[i, 1].bar(np.arange(1,11), prediction)
    axes[i, 1].set_xticks(np.arange(1,11))
    axes[i, 1].set_title("Categorical distribution. Model prediction")

plt.show()

```



```
[86]: num_test_images = x_test.shape[0]

random_inx = np.random.choice(num_test_images, 5)
random_test_images = x_test[random_inx, ...]
random_test_labels = y_test[random_inx, ...]

predictions = model_cnn.predict(random_test_images)

fig, axes = plt.subplots(5, 2, figsize=(16, 12))
fig.subplots_adjust(hspace=0.4, wspace=-0.2)

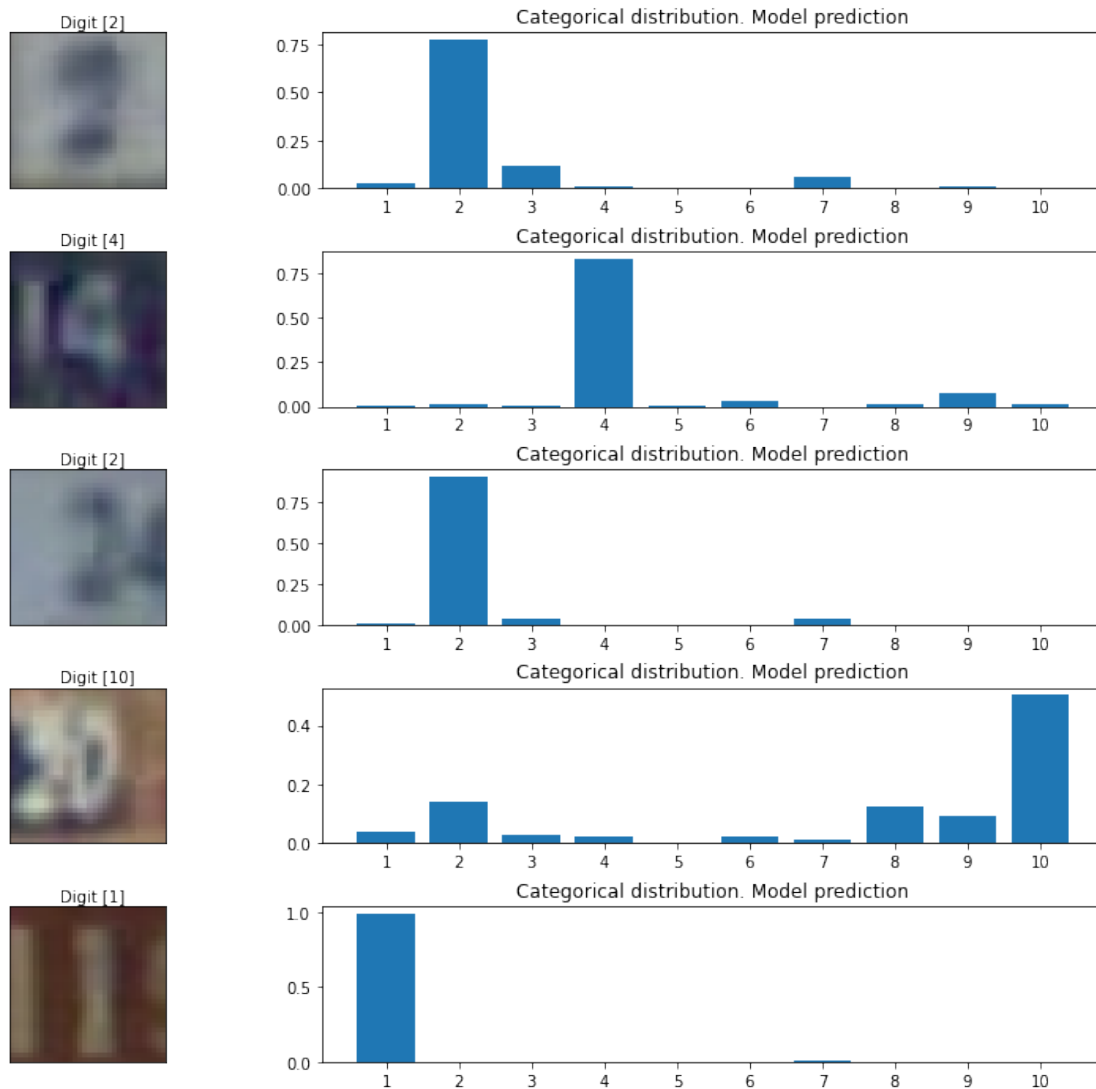
for i, (prediction, image, label) in enumerate(zip(predictions,
    ↳ random_test_images, random_test_labels)):
```

```

axes[i, 0].imshow(np.squeeze(image))
axes[i, 0].get_xaxis().set_visible(False)
axes[i, 0].get_yaxis().set_visible(False)
axes[i, 0].text(10., -1.5, f'Digit {label}')
axes[i, 1].bar(np.arange(1,11), prediction)
axes[i, 1].set_xticks(np.arange(1,11))
axes[i, 1].set_title("Categorical distribution. Model prediction")

```

```
plt.show()
```



```
[ ]:
```

```
[ ]:
```

[ ]: